



MACHINE  
LEARNING  
TOKYO

# Reinforcement Learning

Session - 1

# What is MLT-RL Series?

- Series of study sessions to understand and learn RL.
- Along with application and code for problem statements
- Starting with basic RL understanding
- All the way to Deep RL algorithm understanding and applications
- We try to address them as ["WHY", "WHAT", "HOW"]



MACHINE  
LEARNING  
TOKYO



# Alright!! [ “What” ] is RL?

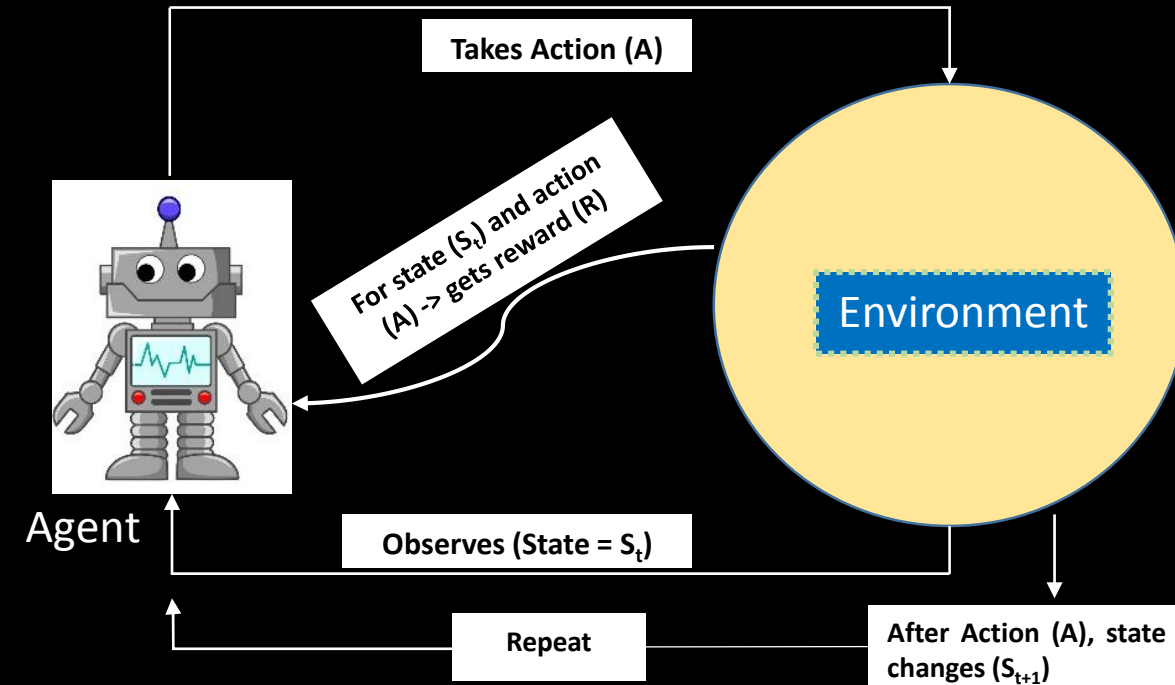
- Proposed originally by Edward L. Thorndike, in principle, Reinforced Learning means, the ability of a system to learn from experiences by taking actions. If actions are linked with rewards, those actions are stamped in memory. He called this “Law of Effects”.

Note:  
Important terminologies specific to a RL problem or mathematical formulation will be marked in “THIS COLOR”.

Edward L. Thorndike Experiment – Video :  
[https://www.youtube.com/watch?time\\_continue=3&v=fanm--WyQJo](https://www.youtube.com/watch?time_continue=3&v=fanm--WyQJo)

# [“What”] Explained

- The objective of the model is
  - Given a “**State ( $S_t$ )**” of the environment
  - Train an “Agent” to take ‘good’ “**Actions ( $A$ )**”
  - In order to maximize the “**Rewards ( $R$ )**” received.
- Conditions
  - “**State ( $S_t$ )**” is not in control of the agent. Agent can only observe the state.
  - “**Reward ( $R$ )**” is also a property of the environment. Observed by the Agent, may be with a time-lag
  - Agent can only control the “**Action ( $A$ )**” that it can take



# [ “What” ] is the Agent trying to achieve?

- Definition of Task

- An RL agent tries to optimize actions to achieve a task given. For example
  - Playing a game
  - Deciding when to “Sell/Buy/Hold” stocks for trading in stock exchange
  - Making a robot walk.
  - Making a baby sleep?? **Really** !! (Duh.. Whatever!!)

- Types of Tasks

- **EPISODIC** : Tasks which end after certain time, for example playing a game. When a game ends (Win/Loss/Draw), the **episode** gets over.
- **CONTINUOUS** : Tasks which (theoretically) do not have an end, for example getting a robot to walk. For such tasks, we divide the continuous tasks in smaller tasks, each being referred to as **episodes**.

# [ “HOW” ] : Parametrization

- Any ML/AI task needs numeric data. Taking this one step ahead, every element of an AI/ML problem needs to be “parametrized” so that it can be represented in the form of numbers
- Therefore, we need to parametrize [ “State”, “Reward”, “Action” ] in an RL problem

# [ “HOW” ] : Understanding “STATE/ACTION/REWARD”



- **State Vector**

- It is the parametrized version of the environment. For eg: While playing Tic-Tac-Toe game, the environment would be an `np.array` of shape of the board (say 3X3).

- **Action**

- Generally, looks like a One-Hot-Encoded-Vector, as the model will suggest one of the actions possible. However, it is important know what actions are possible.

- **Reward**

- **EPISODIC** Tasks : Find seq of tasks that result in majority of episodes successful.
- **CONTINUOUS** Tasks : Avg reward of a time period for small episodes to be maximized.

# [ “HOW” ] : State/Action/Reward Examples

- Investment Portfolio (CONTINUOUS Task)

**State** :  $f(s_t) = (\text{Cash} - \text{in} - \text{hand}, \% \text{ investment}, \text{liabilities}, \text{assets}, \text{etc.})$

**Action** : Sell/Buy/Hold stock

**Reward** : Avg monthly Profit/Loss (%)

---

- Cab Driver Ride Selection (CONTINUOUS Task)

**State** :  $f(s_t) = (\text{Source Location}, \text{Target location}, \text{Fuel available}, \text{time}, \text{day}, )$

**Action** : Accept Ride / Reject Ride

**Reward** : Avg Monthly earning

---

- Tic-Tac-Toe Game (EPISODIC Task)


**State** :  $f(s_t) = \text{np.array} \rightarrow [1,0,\text{np.nan}]$  based on [ “X”, ”O”, free space] on the board (respectively)

**Action** : Location to play the move. (Play “X” at [0,2] location)

**Reward** : Game Won/Loss/Draw.



# [ “HOW” ] : Markovian Assumption




## SCENARIO

- Say you are at the SOURCE and want to reach the DESTINATION. And there are some paths that you can follow
- To know, if taking LEFT is better or RIGHT, when at SOURCE, you would really want to know, what lies ahead at each node and see in totality which action would be better.
- In real life, such information may not be available.

# [ “HOW” ] : Markovian Assumption

## MARKOV ASSUMPTION



- DEFINITION

- Given the current STATE and ACTION TAKEN, the FUTURE can be predicted. It is independent of what was in past.

- EXPLANATION



- ACTION to be taken in current STATE is dependent only on the current STATE.
- To predict a good ACTION, the current STATE vector has all required information
- A good ACTION is not dependent on how we reached the current state and from what path.

**Also known as “One Markovian Assumption”.**

# [ “HOW” ] : Markov Decision Process (MDP)

- MDP means, the problem statement at hand, can be described in *STATE TRANSITION MANNER*, based on some *PROBABILISTIC* conditions.
- In MDP, the environment and AGENT both follow Markovian Assumption
- Decision = Action taken by the AGENT is based on Markovian Assumption
- Markov = It is symbolic of the STATE vector, which is assumed to provide all relevant information for taking the action & predicting the future.
- MDP assumes that every STATE in that process is dependent only on the previous state
- ***Not all ML problems can be formulated in an MDP***
- If a problem can be formulated in an MDP format i.e. (STATE, ACTION, REWARD) structure only then it is good candidate for application of RL.

# [ ``HOW`` ] : Model of environment & Types of RL Algorithms





- DEFINITION :

- The probabilistic information pertaining to the state transitions in an environment when an action is taken.

- EXPLANATION :

- When in STATE(S1), if agent takes a LEFT, then what is the probability that it will reach STATE(S2). Based on above it is 100%
- However, in some cases, taking the same action in a particular state may lead to different outcomes. For example : A healthy person is bit by a mosquito, he may or may not fall ill. This becomes probabilistic in nature.

# [ ``HOW`` ] : Model of environment & Types of RL Algorithms





- FORMULATION

$$\textit{Model of Environment} = p(s', r \mid s, a)$$

- Probability of reaching a STATE ( $s'$ ) and getting a REWARD ( $r$ ) given that in current STATE ( $s$ ), an action ( $a$ ) was taken.

# [ ``HOW`` ] : Model of environment & Types of RL Algorithms



- RL Algorithm Types

- **Model Based Methods Algorithms**

These algorithms assume that such a probabilistic map  $p(s', r | s, a)$  is available for modelling.

- **Model Free Methods Algorithms**

These models do not assume the availability of probabilistic map, they are based on the theory of experimentation and learning along the way.

[ “HOW” ] : All said, what does the agent learn

- Given a particular **STATE** of the environment, the job of the AGENT is to predict a good **ACTION**
- Making the **AGENT** learn this ability is the main objective of RL modelling.
- In the language of RL, this objective is called **POLICY** or **CONTROL OBJECTIVE**.

[ “HOW” ] : All said, what does the agent learn

## Food for Thoughts:

- Think about what a **POLICY** should look like, what information should it capture?
- For example
  - Given a **STATE**, what is the best **ACTION** (main objective) ?
  - Given a **STATE**, differentiate empirically that some **ACTIONS** are better than other
- These 2 elements are key for taking a good decision
- So.....



[ “HOW” ] : All said, what does the agent learn

- The Policy looks something like this

$$\textit{Policy} = \pi ( \textit{action} \mid \textit{state} )$$

- Probability of taking an ACTION, given the STATE
- This probability map is what the agent aims to learn, so as to predict one of the possible actions

# Quick Recap : What we know....


- We know, what are **STATE**, **ACTION** & **REWARDS** in an RL Problem.
- What are **EPISODIC** & **CONTINUOUS** tasks for an RL agent.
- What is parametrization of **STATE**, **ACTION** & **REWARDS**
- We know, what is **Markovian Assumption**.
- We know, what is **Markov Decision Process (MDP)**
- We know, what is model of environment (Model Based/Model Free Methods)
- We know, what is **POLICY** with respect to an AGENT and in terms of an RL Problem.

# [ “HOW” ] : RL Equations



MACHINE  
LEARNING  
TOKYO

- Say you build an RL agent to play the game of CounterStrike.
- In the scene below, not all places/positions in the space visible would be considered as “good” for a player.



- Human logic suggests that State (S1) – “Ducking” is a better than State (S2) – “Standing”. A less visible place, with maximum attack range.
- Therefore, some states are good, some are “not so good”.

# [ “HOW” ] : RL Equations

- How to decide, which **STATE** is better
  - In the CounterStrike example, may be, one way to judge that is, how long the player is “alive” in the game by being in that **STATE**. Longer the player stays alive and also with more targets, the better the state.
  - Inturn, what is the *immediate reward* & *cumulative future rewards*.
- How to decide which **ACTION** is better
  - Again in the Counter Strike example, if the player is under attack, would “ducking” be a better action, or “standing” and shooting be a better action. Again, the answer lies in measuring the rewards for it.

## BIG QUESTION

- How does an AGENT possibly learn something like this??
- What Equation?

# [ “HOW” ] : RL Equations & Algorithms



MACHINE  
LEARNING  
TOKYO



**OVER TO THE BOARD!!**



# In brief, now we know...

- What are **STATE**, **ACTION** & **REWARDS** in an RL Problem.
- What are **EPISODIC** & **CONTINUOUS** tasks for an RL agent.
- What is parametrization of **STATE**, **ACTION** & **REWARDS**
- What is **Markovian Assumption** & **Markov Decision Process (MDP)**
- What is model of environment (Model Based/Model Free Methods).
- What is **POLICY** with respect to an AGENT and in terms of an RL Problem.
- What are **STATE VALUE FUNCTION**  $v(s)$  & **ACTION VALUE FUNCTION**  $q(a | s)$
- For Model Based Methods, what is **Value Iteration algorithms** & what is **Policy Iteration algorithms**

$$\textit{Policy} = \pi ( \textit{action} \mid \textit{state} )$$

$$\textit{Model of Environment} = p(s', r \mid s, a)$$

$$\textit{State Value } v_{\pi}(s) = \sum_a \pi (a \mid s) q_{\pi}(s, a)$$

$$\textit{Action Value } q_{\pi}(s, a) = \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')]$$

Combined equation for State Value & Action Value

$$\textit{State Value } v_{\pi}(s) = \sum_a \pi (a \mid s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma v(s')]$$

(s) - Current State

(a) - Action

(s') - Next State

(r) - Immediate Reward

( $\pi$ ) - Policy

( $\gamma$ ) - Discount Factor



MACHINE  
LEARNING  
TOKYO

THANK YOU!

We'll be back with **Session -2** soon!!